Challenges of Sparse Bayesian Inversion and Uncertainty Quantification

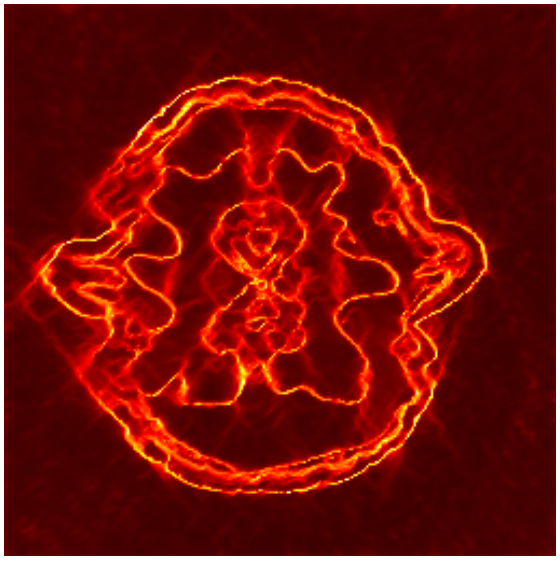
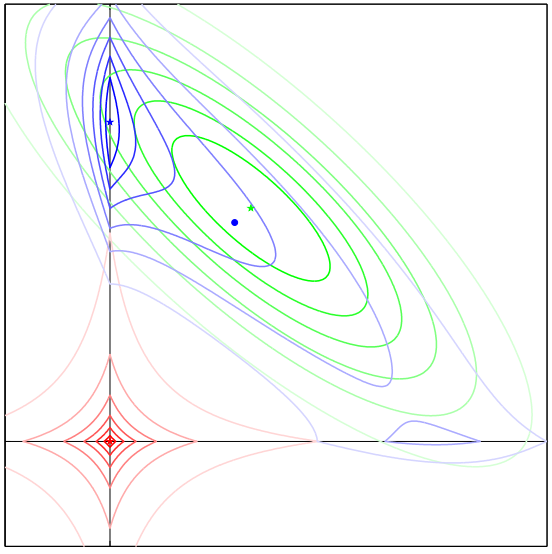
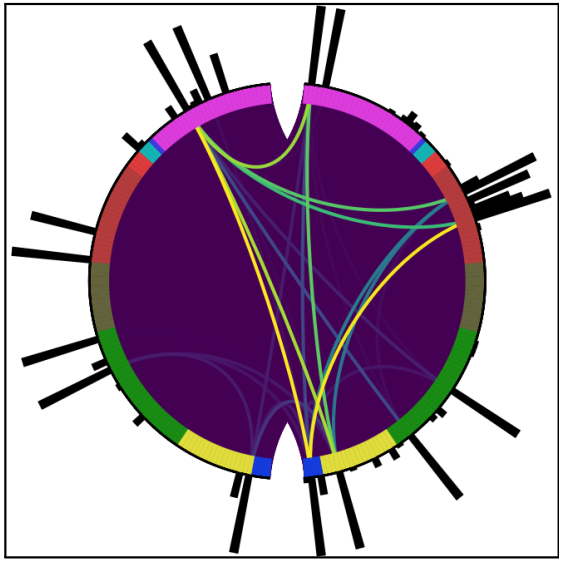

Felix Lucka
University College London
f.lucka@ulc.ac.uk

cmic
Centre for Medical Image Computing

Bayesian & Nonlinear Inverse Problems
Lorenz Center, Aug 29, 2017.

UCL CENTRE FOR INVERSE PROBLEMS

Traditional task: Produce results to be interpreted by trained experts

$\Longrightarrow$ *Qualitative* usage of the reconstructed information.

Example: Conventional *computer tomography* (*CT*).

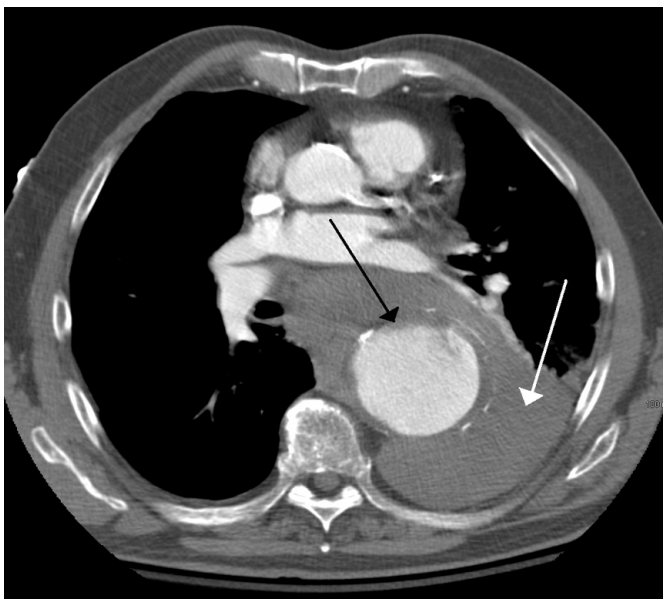

Source: Wikimedia Commons

Traditional task: Produce results to be interpreted by trained experts

$\Longrightarrow$ *Qualitative* usage of the reconstructed information.

New demand: Produce results for automatized analysis procedures / hypothesis testing; multimodal imaging.

$\Longrightarrow$ *Quantitative* usage of the reconstructed information.

Example: *Dynamical causal modeling* (*DCM*).

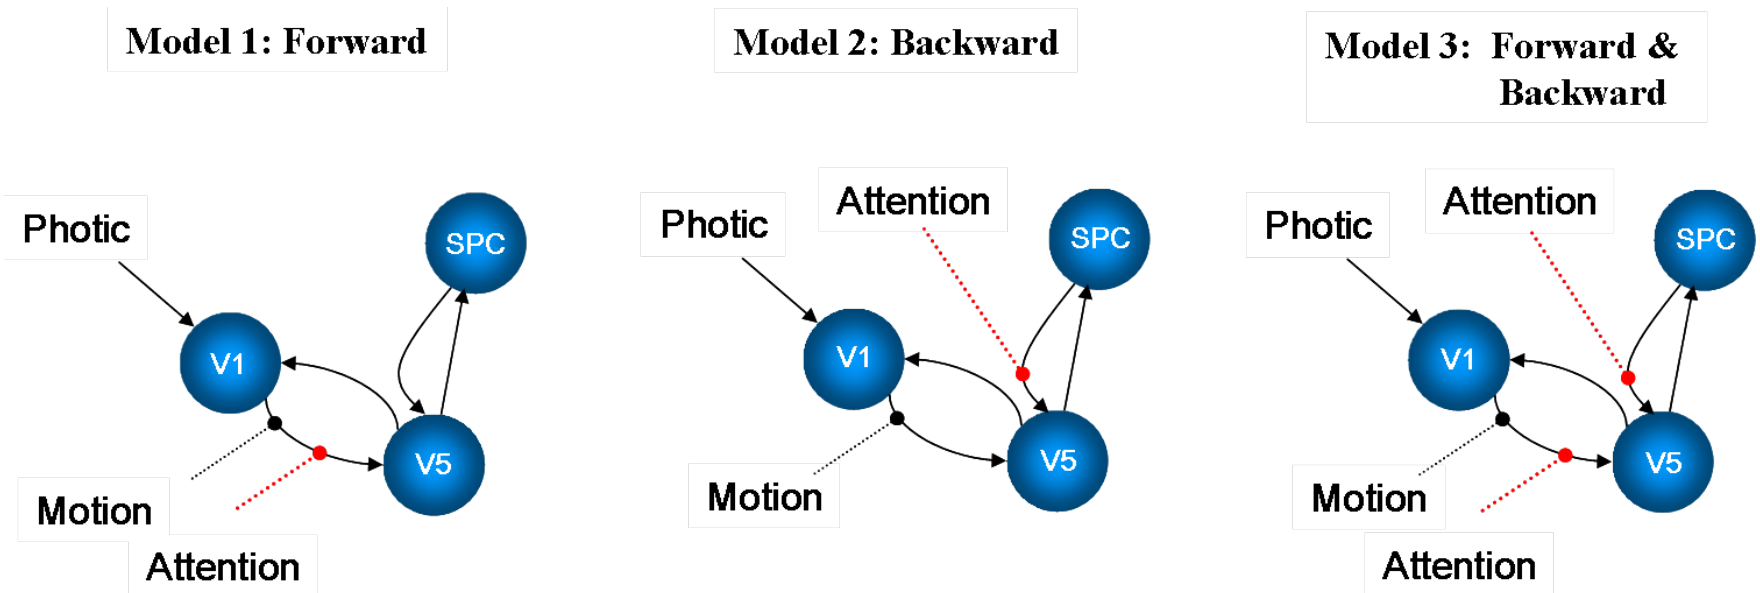

Source: Andre C. Marreiros et al. (2010), Scholarpedia, 5(7):9568.

Noisy, ill-posed inverse problems:

$$f = N(\mathcal{A}(u), \varepsilon)$$

Example: $f = Au + \varepsilon$

$p_{like}(f|u) \propto$
$\exp\left(-\frac{1}{2}\|f - A\,u\|_2^2\right)$

$p_{prior}(u) \propto$
$\exp\left(-\lambda\,\|D^T u\|_2^2\right)$

$p_{post}(u|f) \propto$
$\exp\left(-\frac{1}{2}\|f - A\,u\|_2^2 - \lambda\,\|D^T u\|_2^2\right)$

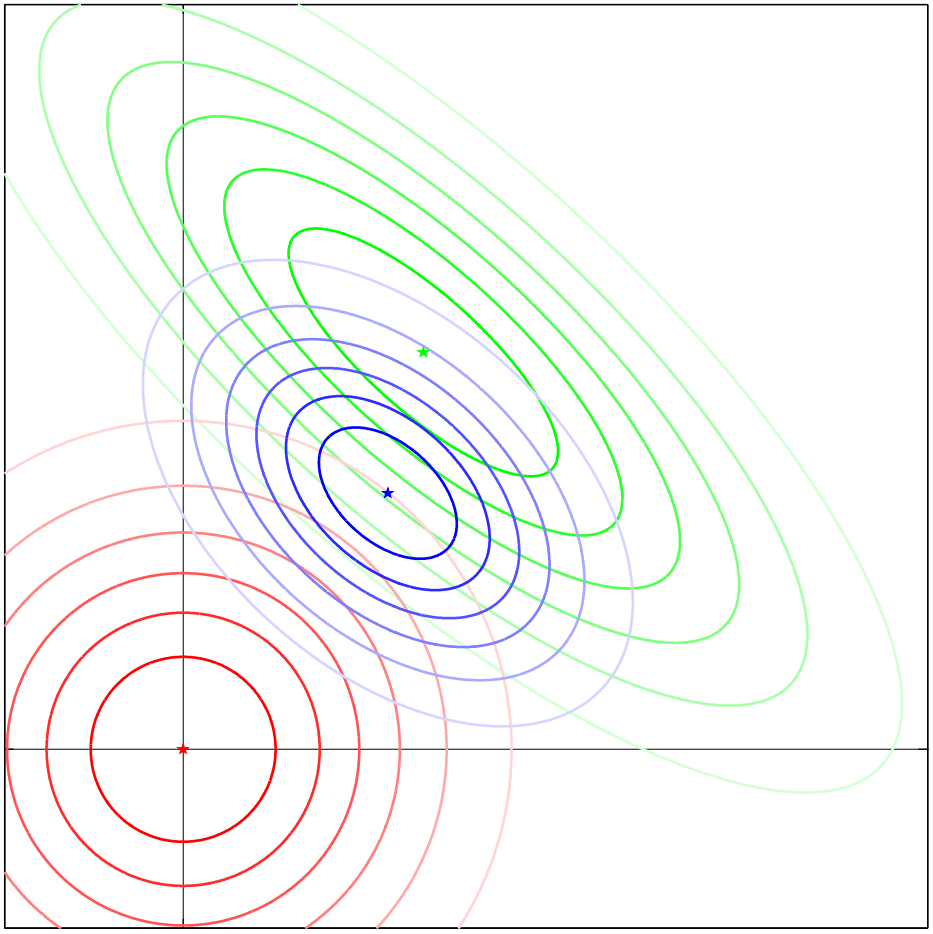

Probabilistic representation allows for a rigorous quantification of the solution's uncertainties.

Noisy, ill-posed inverse problems:

$$f = N(\mathcal{A}(u), \varepsilon)$$

Example: $f = Au + \varepsilon$

$p_{like}(f|u) \propto$
$\exp\left(-\frac{1}{2}\|f - Au\|_2^2\right)$

$p_{prior}(u) \propto$
$\exp\left(-\lambda \|D^T u\|_1\right)$

$p_{post}(u|f) \propto$
$\exp\left(-\frac{1}{2}\|f - Au\|_2^2 - \lambda \|D^T u\|_1\right)$

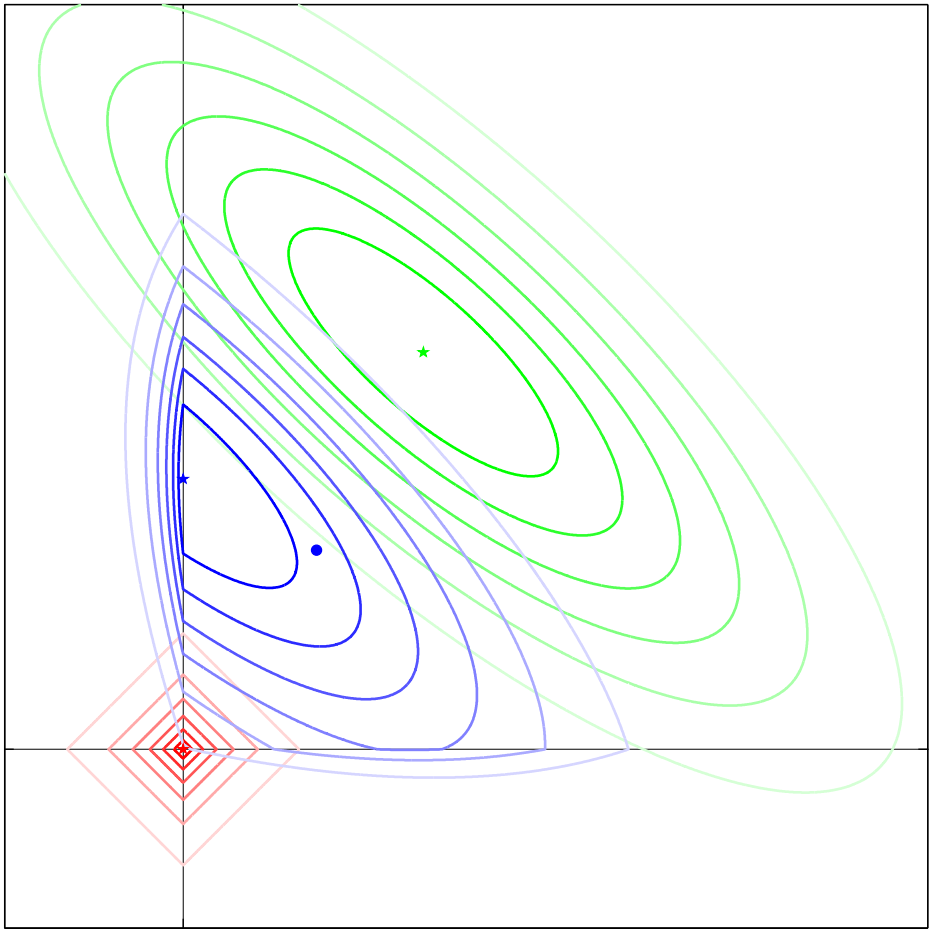

Probabilistic representation allows for a rigorous quantification of the solution's uncertainties.

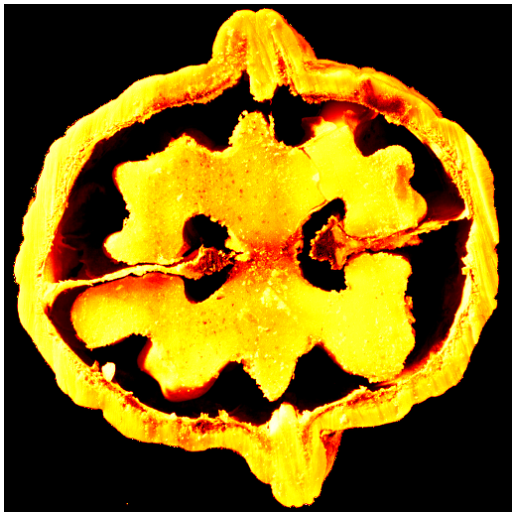

(a) 100%

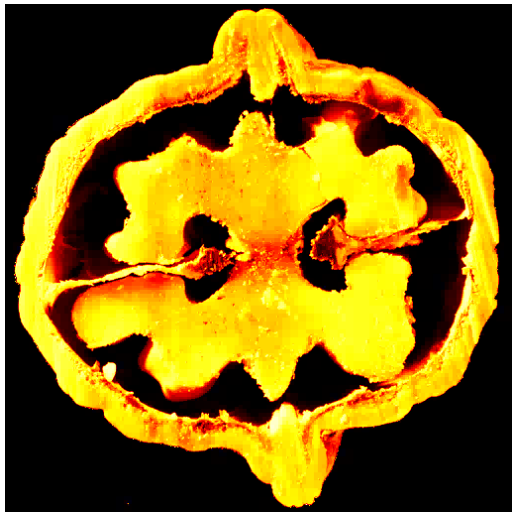

(b) 10%

(c) 1%

Sparsity a-priori constraints are used in variational regularization, compressed sensing and variable selection:

$$\hat{u}_\lambda \quad = \underset{u}{\operatorname{argmin}} \left\{ \tfrac{1}{2} \| f - A\,u \|_2^2 + \lambda \| D^T u \|_1 \right\}$$

(e.g. *total variation*, *wavelet shrinkage*, *LASSO*,...)

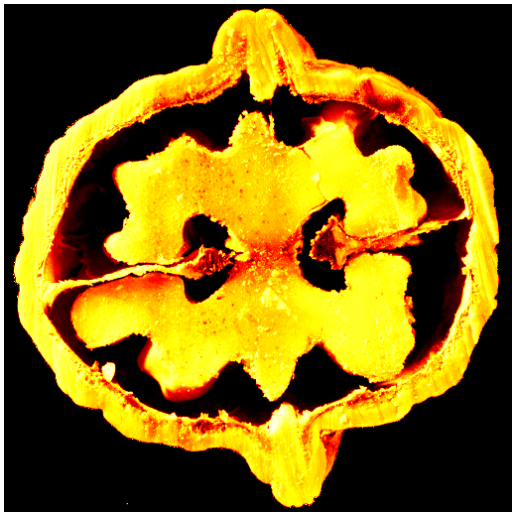

(a) 100%

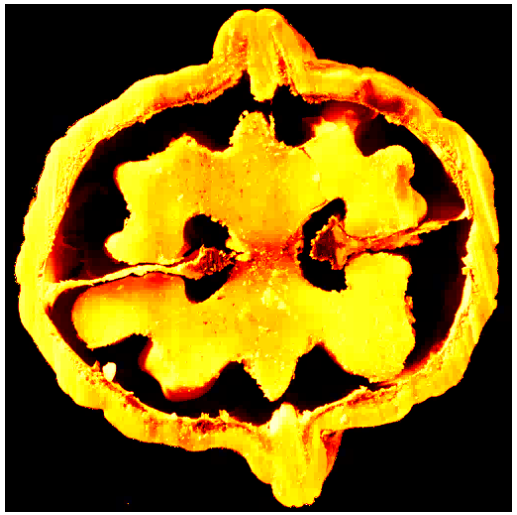

(b) 10%

(c) 1%

Sparsity a-priori constraints are used in variational regularization, compressed sensing and variable selection:

$$\hat{u}_\lambda \quad = \underset{u}{\operatorname{argmin}} \left\{ \tfrac{1}{2} \| f - A\,u \|_2^2 + \lambda \| D^T u \|_1 \right\}$$

(e.g. *total variation*, *wavelet shrinkage*, *LASSO*,...)

Sparse Bayesian inversion?

Uncertainty Quantification for Sparse Bayesian Inversion

UCL

- How to model sparsity?
  - ℓ_1-norm priors.
  - Gaussian scale mixture (hierarchical Bayesian)
  - ℓ_p-norm scale mixture (hierarchical Bayesian)
- How to we compute estimators / UQ measures?
- What can we say about estimators?
- Meaningful UQ measures for sparse inversion / imaging?

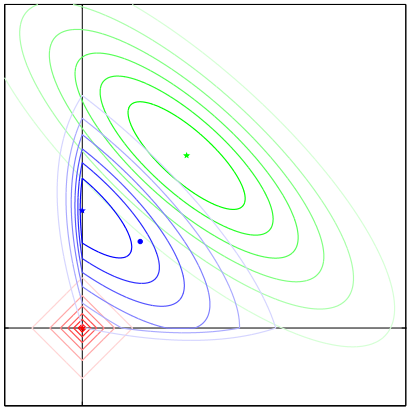

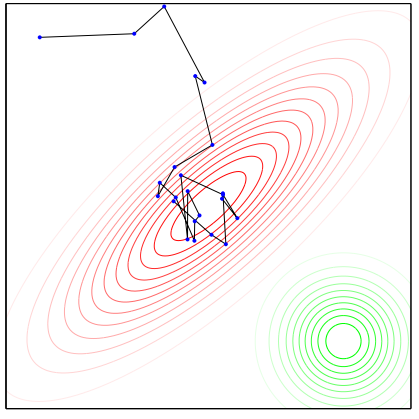

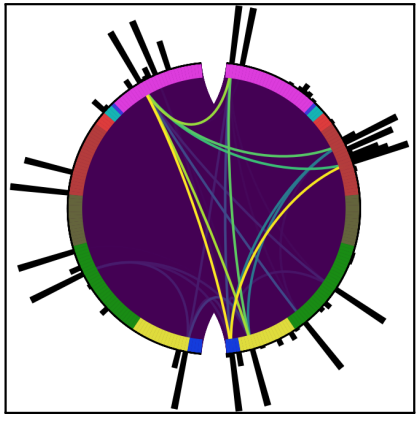

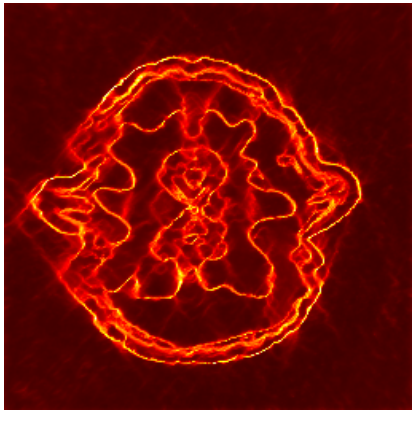

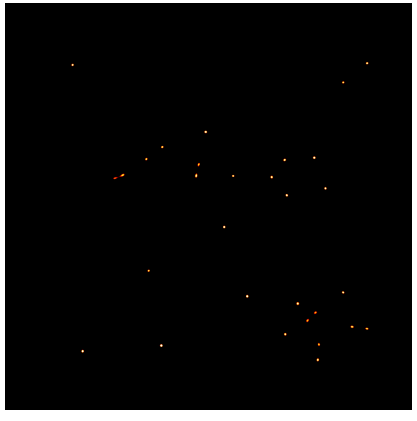

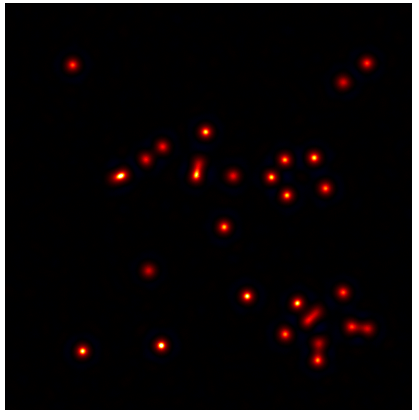

Efficient MCMC for Sparse Image Reconstruction

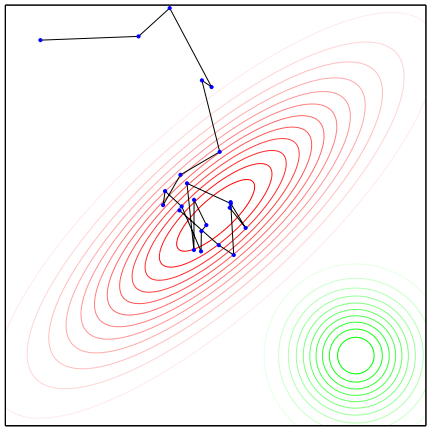

Task: Monte Carlo integration by samples from

$$p_{post}(u|f) \propto \exp\left(-\tfrac{1}{2}\|f - A\,u\|^2_{\Sigma_\varepsilon^{-1}} - \lambda\,\|D(u)\|_1\right)$$

Problem: Standard Markov chain Monte Carlo (MCMC) sampler (Metropolis-Hastings) inefficient for large n or λ.

Efficient MCMC for Sparse Image Reconstruction

Task: Monte Carlo integration by samples from

$$p_{post}(u|f) \propto \exp\left(-\tfrac{1}{2}\|f - A\,u\|^2_{\Sigma_\varepsilon^{-1}} - \lambda\,\|D(u)\|_1\right)$$

Problem: Standard Markov chain Monte Carlo (MCMC) sampler (Metropolis-Hastings) inefficient for large n or λ.

Contributions:

- Development of different Gibbs sampler implementations.
- Still efficient for high-dimensional imaging ($n > 10^6$).

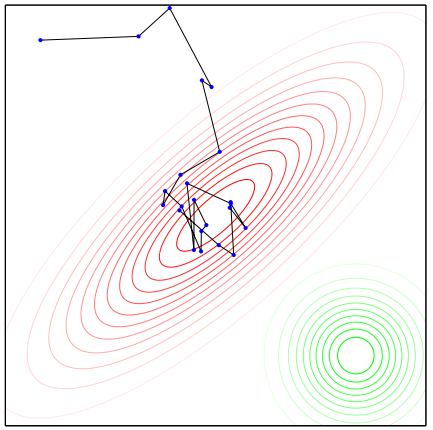

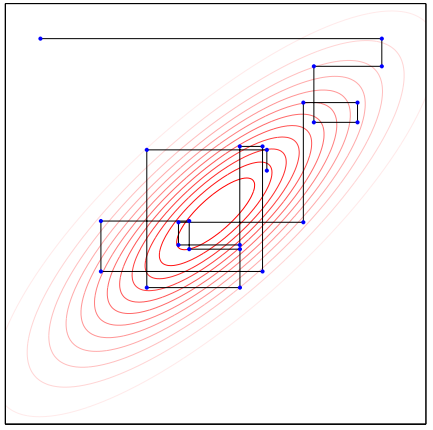

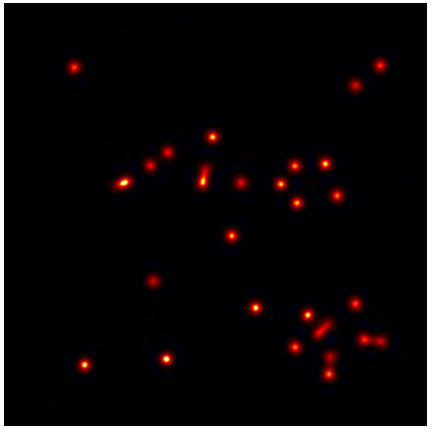

F.L, 2016. *Fast Gibbs sampling for high-dimensional Bayesian inversion*, Inverse Problems.

F.L, 2012. *Fast Markov chain Monte Carlo sampling for sparse Bayesian inference in high-dimensional inverse problems using L1-type priors*, Inverse Problems.

Efficient MCMC for Sparse Image Reconstruction

Task: Monte Carlo integration by samples from

$$p_{post}(u|f) \propto \exp\left(-\tfrac{1}{2}\|f - A\,u\|^2_{\Sigma_\varepsilon^{-1}} - \lambda\,\|D(u)\|_1\right)$$

Problem: Standard Markov chain Monte Carlo (MCMC) sampler (Metropolis-Hastings) inefficient for large n or λ.

Work by Marcelo Pereyra et al.:

- Unadjusted Langevin algorithm applied to Moreau-Yoshida envelopes of posterior energy.
- As easy to implement as proximal gradient descent.

Alain Durmus, Eric Moulines, Marcelo Pereyra, 2016. *Efficient Bayesian computation by proximal Markov chain Monte Carlo: when Langevin meets Moreau*, arXiv:1612.07471.

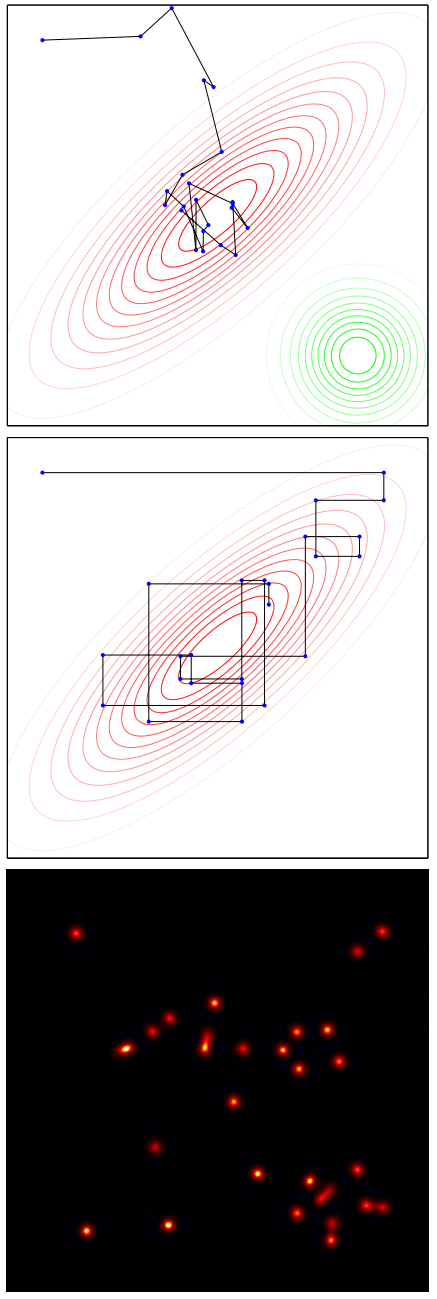

New Theoretical Ideas for an Old Bayesian Debate

$$\hat{u}_{\mathrm{MAP}} := \operatorname*{argmax}_{u \in \mathbb{R}^n} \{\, p_{post}(u|f)\} \quad \text{vs.} \quad \hat{u}_{\mathrm{CM}} := \int u\, p_{post}(u|f)\, \mathrm{d}u$$

- CM preferred in theory, dismissed in practice.
- MAP discredited by theory, chosen in practice.

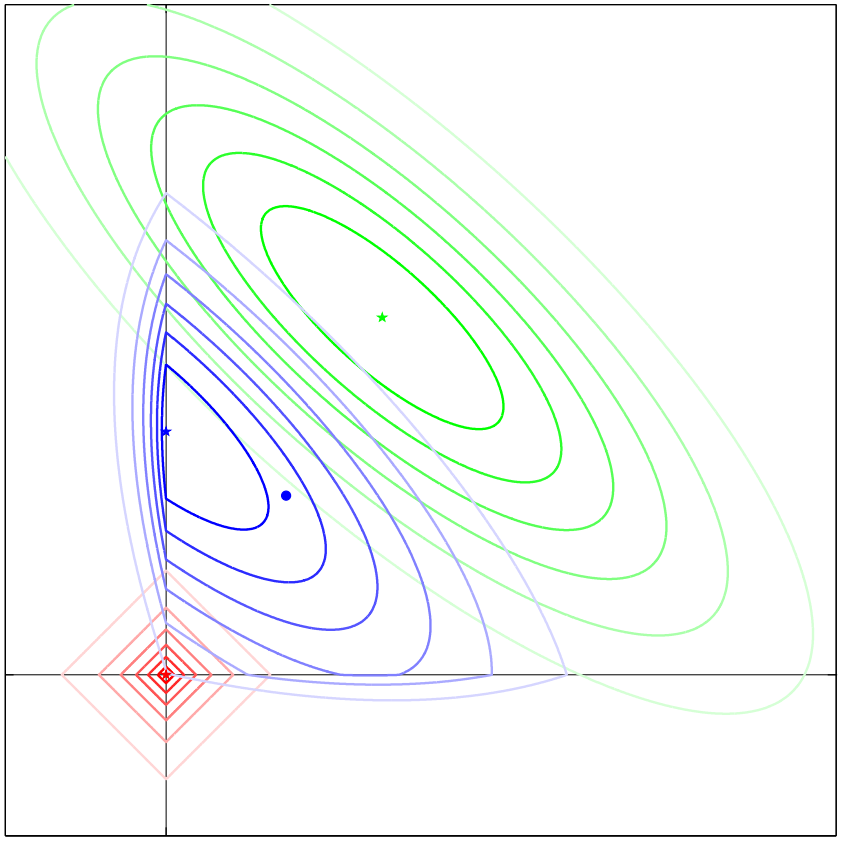

New Theoretical Ideas for an Old Bayesian Debate

$$\hat{u}_{\text{MAP}} := \operatorname*{argmax}_{u \in \mathbb{R}^n} \{\, p_{post}(u|f) \} \quad \text{vs.} \quad \hat{u}_{\text{CM}} := \int u \; p_{post}(u|f) \, \mathrm{d}u$$

- CM preferred in theory, dismissed in practice.
- MAP discredited by theory, chosen in practice.

However:

- MAP results looks/performs better or similar to CM.
- Gaussian priors: MAP = CM. Funny coincidence?
- Theoretical argument has a logical flaw.

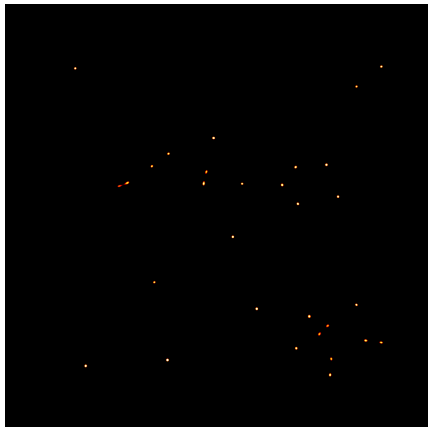

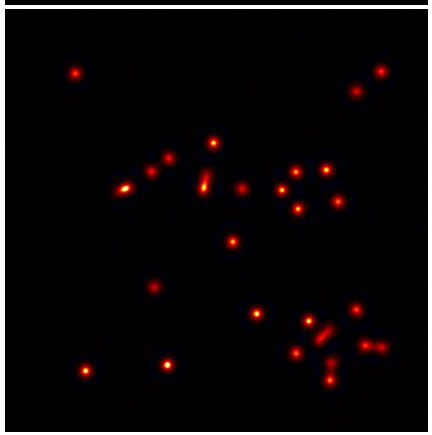

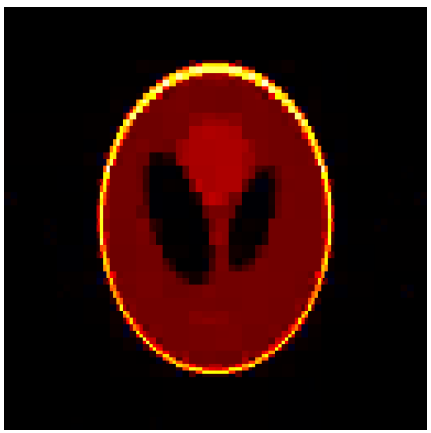

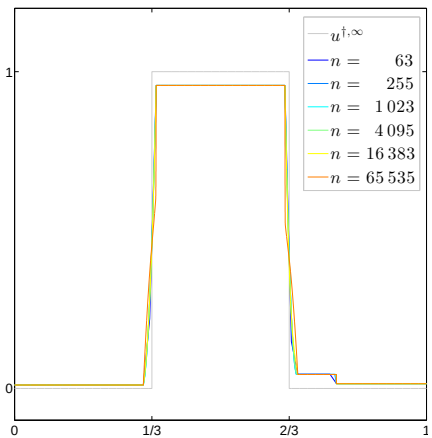

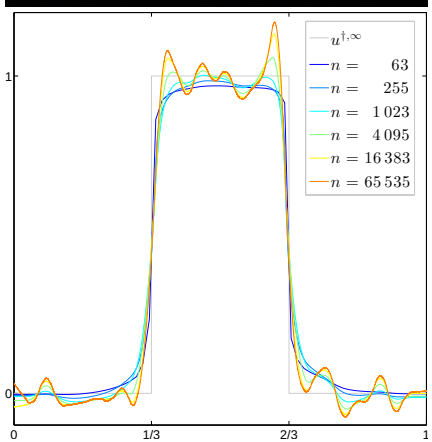

New Theoretical Ideas for an Old Bayesian Debate

$$\hat{u}_{\mathrm{MAP}} := \operatorname*{argmax}_{u \in \mathbb{R}^n} \{\, p_{post}(u|f) \} \quad \text{vs.} \quad \hat{u}_{\mathrm{CM}} := \int u \; p_{post}(u|f) \, \mathrm{d}u$$

- CM preferred in theory, dismissed in practice.
- MAP discredited by theory, chosen in practice.

Contributions:

- Theoretical rehabilitation of MAP.
- Key: Bayes cost functions based on Bregman distances.
- Gaussian case consistent in this framework.

Burger & L, 2014. *Maximum a posteriori estimates in linear inverse problems with log-concave priors are proper Bayes estimators*, *Inverse Problems*, 30(11):114004.

Helin & Burger, 2015. *Maximum a posteriori probability estimates in infinite-dimensional Bayesian inverse problems*, *Inverse Problems*, 31(8)

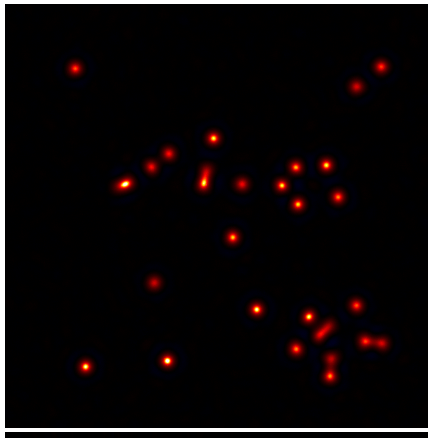

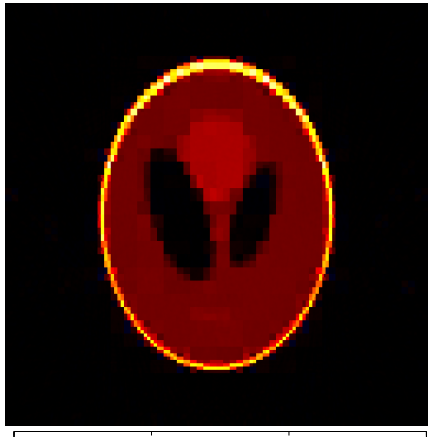

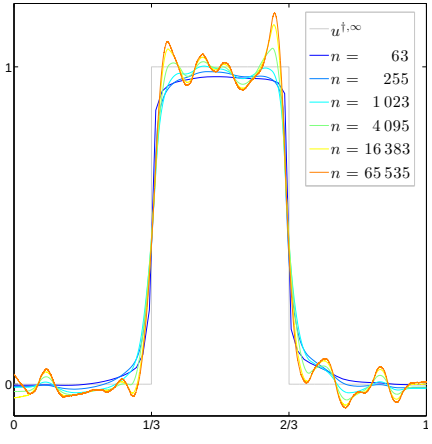

Experimental Data: Limited-Angle CT

- Cooperation with Samuli Siltanen, Esa Niemi et al.
- Besov and TV prior; non-negativity constraints.
- Stochastic noise modeling.
- Uncertainty quantification for limited angle CT.


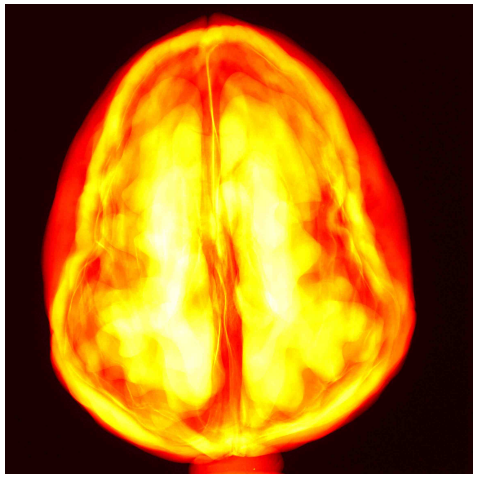
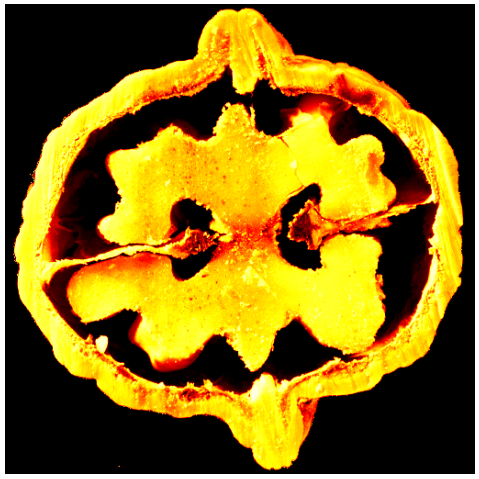
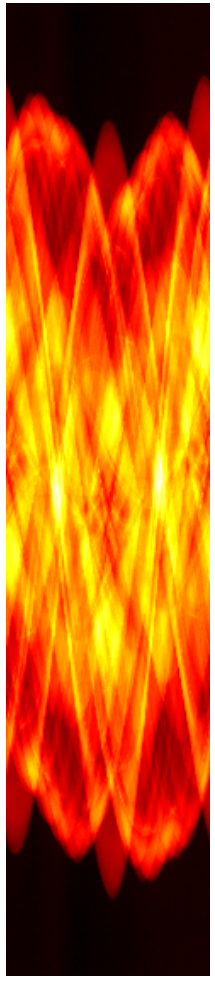

Use the data set for your own work:
http://www.fips.fi/dataset.php (documentation: arXiv:1502.04064)

Walnut-CT with TV Prior: Full vs. Limited Angle

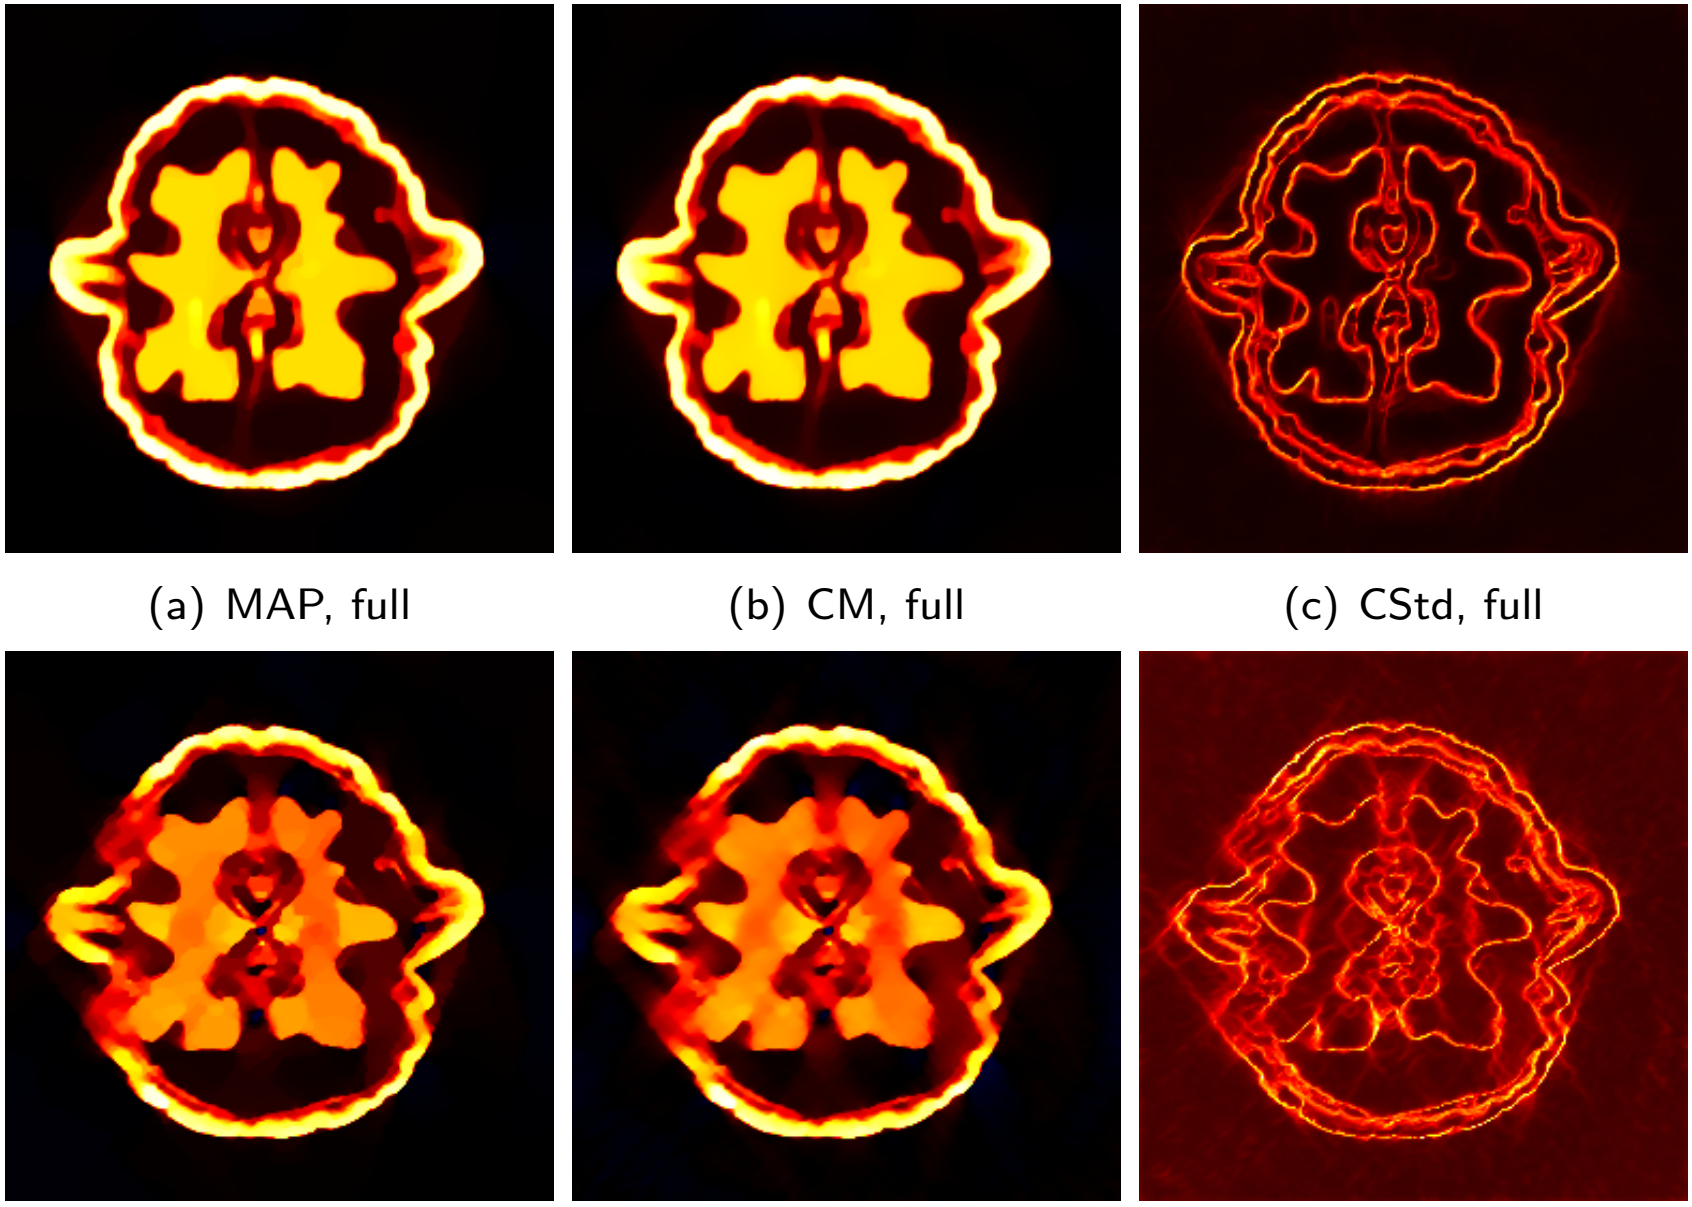

(a) MAP, full (b) CM, full (c) CStd, full

(d) MAP, limited (e) CM, limited (f) CStd, limited

Walnut-CT with TV Prior: Non-Negativity Constraints, Limited Angle

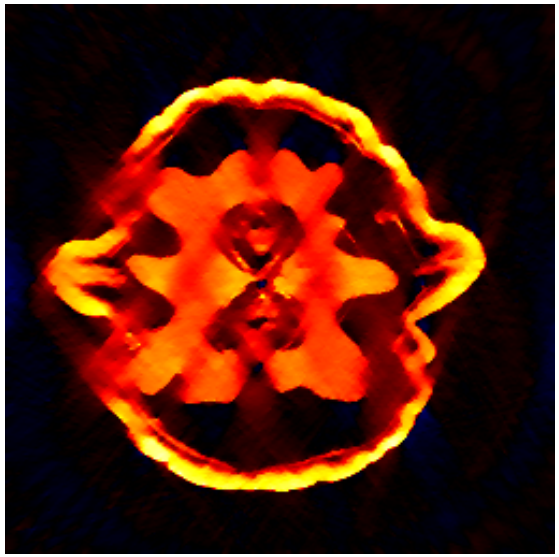

(a) CM, uncon

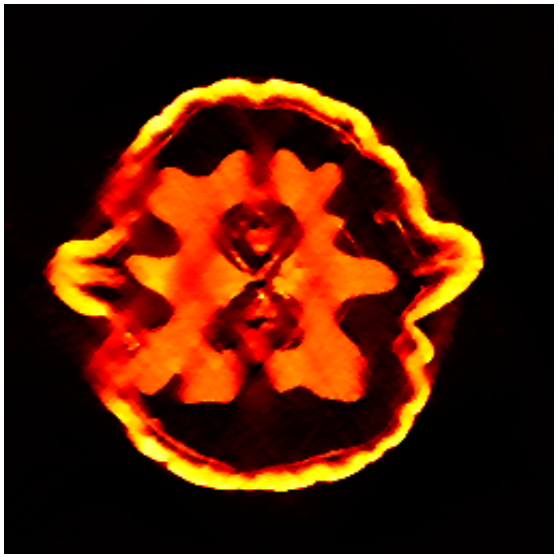

(b) CM, non-neg

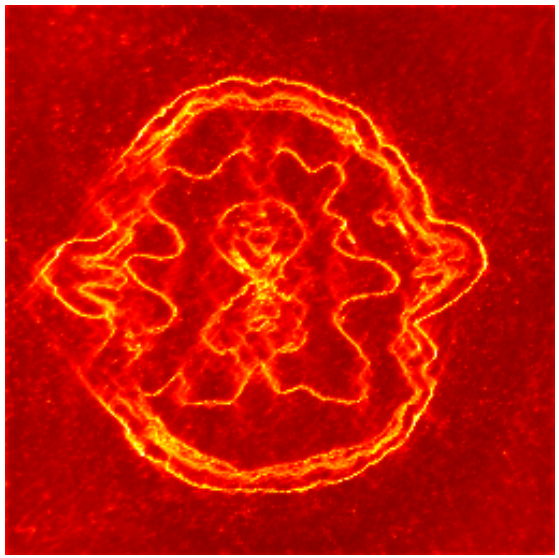

(c) CStd, uncon

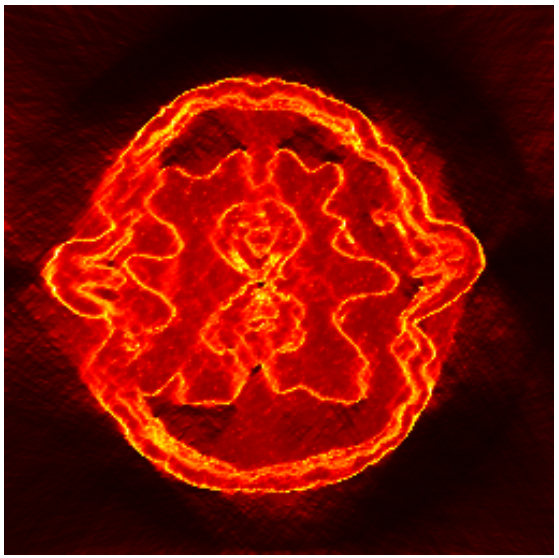

(d) CStd, non-neg

However...

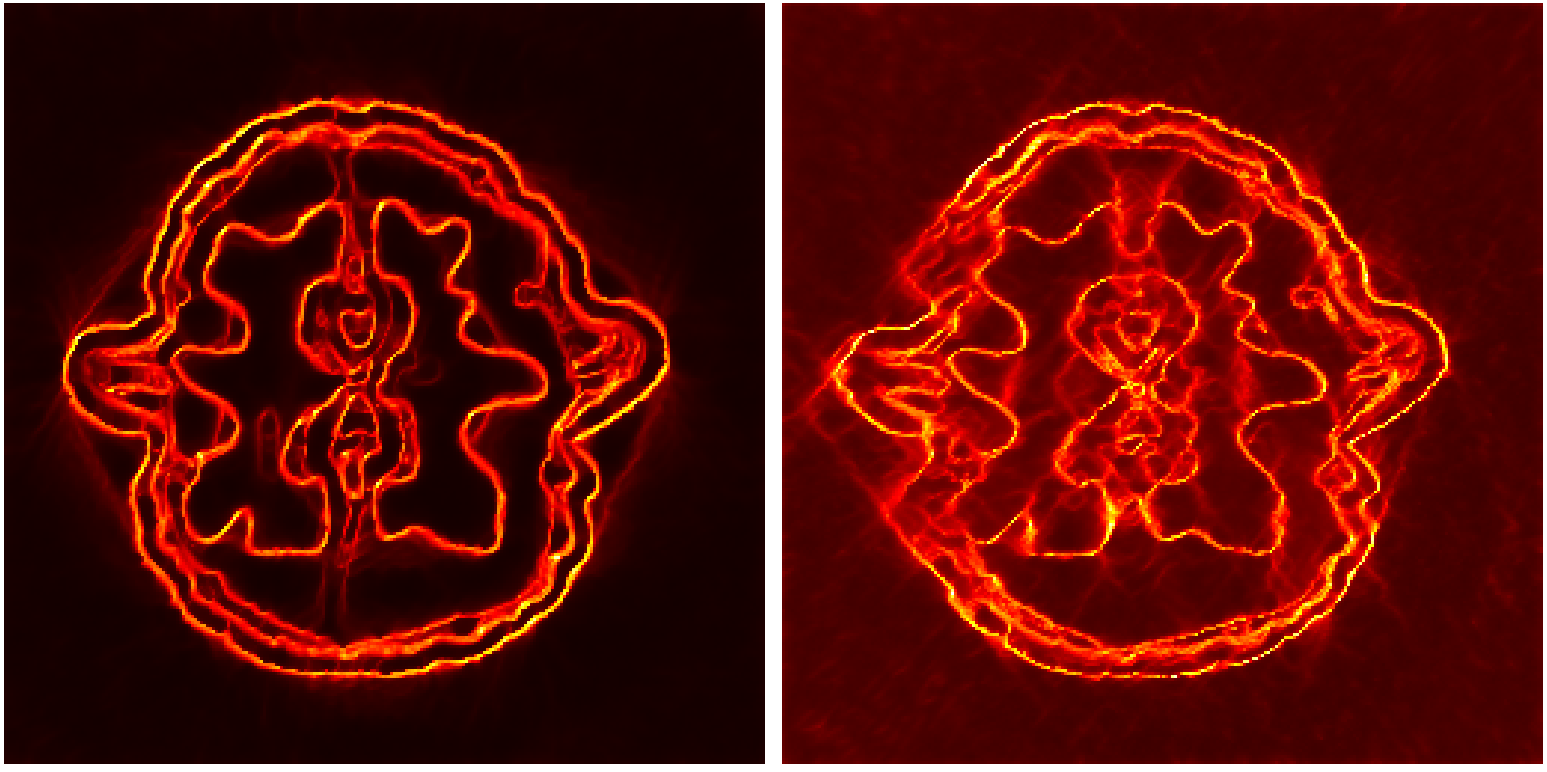

(a) CStd, full

(b) CStd, limited

- What does it really tell me?
- Why does the uncertainty decrease?!

Gaussian increment prior:

$$p_{prior}(u) \propto \prod_i \exp\left(-\frac{(u_{i+1}-u_i)^2}{\gamma}\right)$$

- Gaussian variables take values on a characteristic scale, determined by γ.
- Similar amplitudes are likely, sparsity (= outliers) is unlikely.

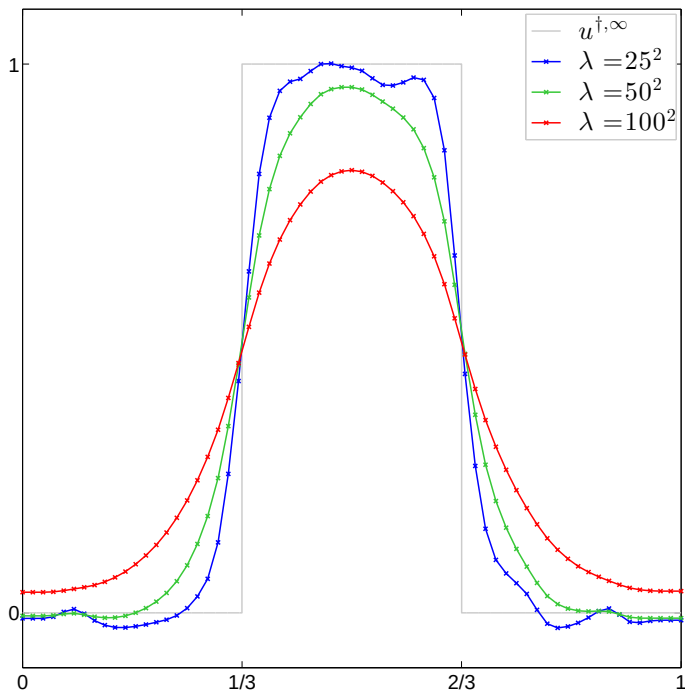

Conditionally Gaussian increment prior:

$$p_{prior}(u|\gamma) \propto \prod_i \exp\left(-\frac{(u_{i+1}-u_i)^2}{\gamma_i}\right)$$

Scale-invariant hyperprior to approximate un-informative γ_i^{-1} prior:

$$p_{hyper}(\gamma_i) \propto \gamma_i^{-(\alpha+1)} \exp\left(-\frac{\beta}{\gamma_i}\right), \quad \textit{inverse gamma distribution}$$

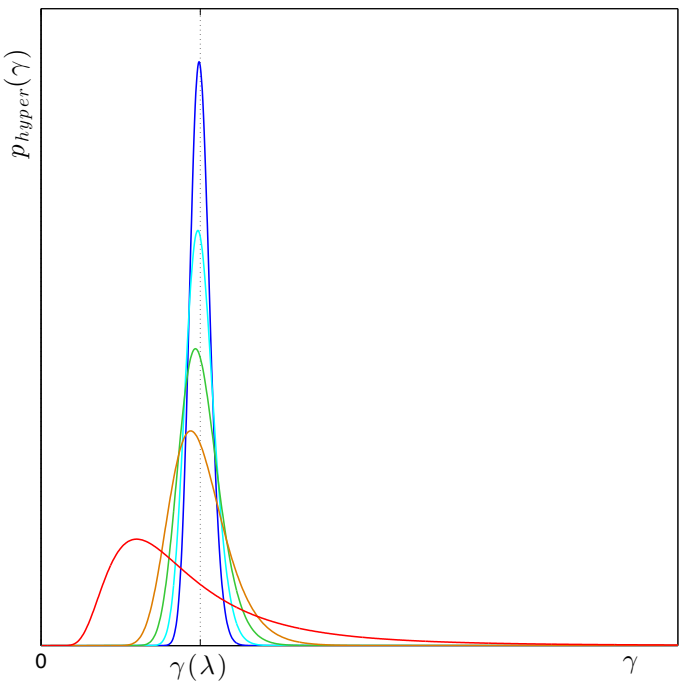

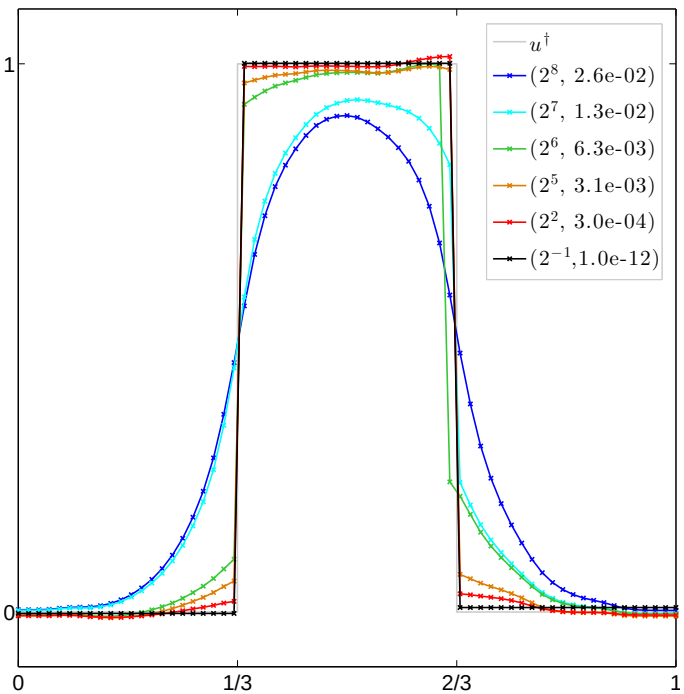

The Implicit Energy Functional behind HBM

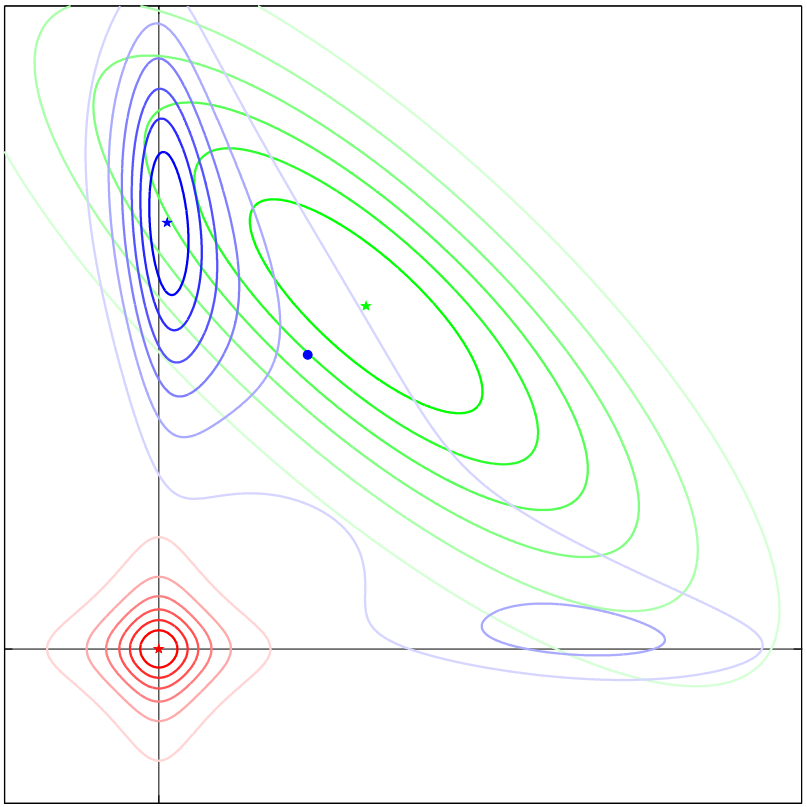

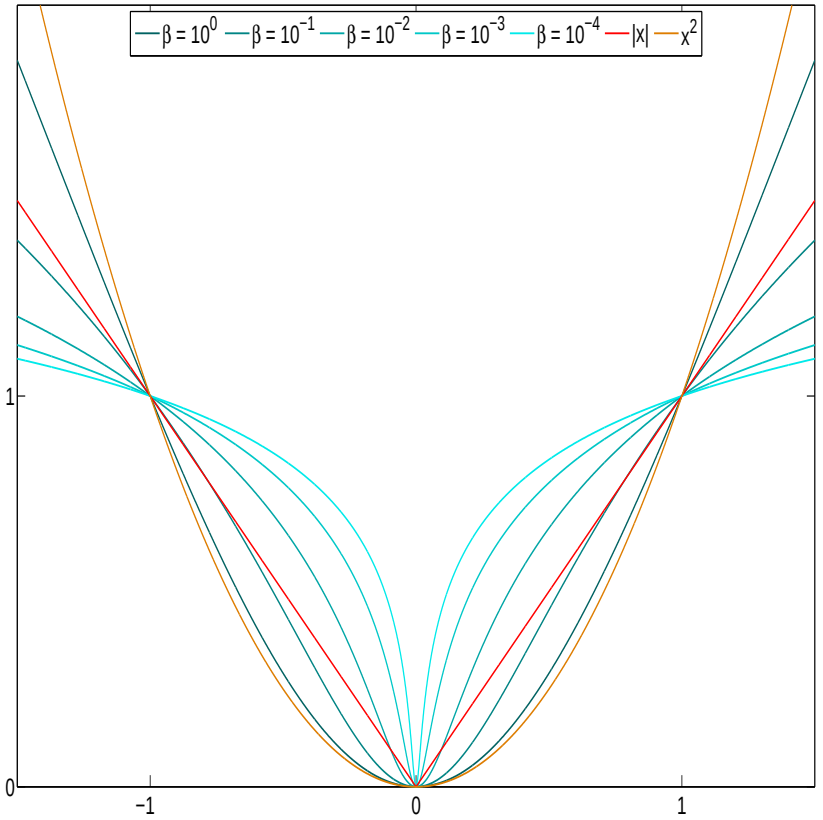

Implicit prior is a Student's t-prior with $\nu = 2\alpha, \theta = \beta/(2\alpha)$:

$$p_{prior}(u) \propto \prod_i \left(1 + \frac{u_i^2}{\nu\theta}\right)^{-\frac{\nu-1}{2}}$$

$$p_{post}(u|f) \propto \exp\left(-\tfrac{1}{2}\|f - A\,u\|^2_{\Sigma_\varepsilon^{-1}} - \frac{\nu-1}{2}\sum_i \log\left(1 + \frac{u_i^2}{\nu\theta}\right)\right)$$

Prior Samples

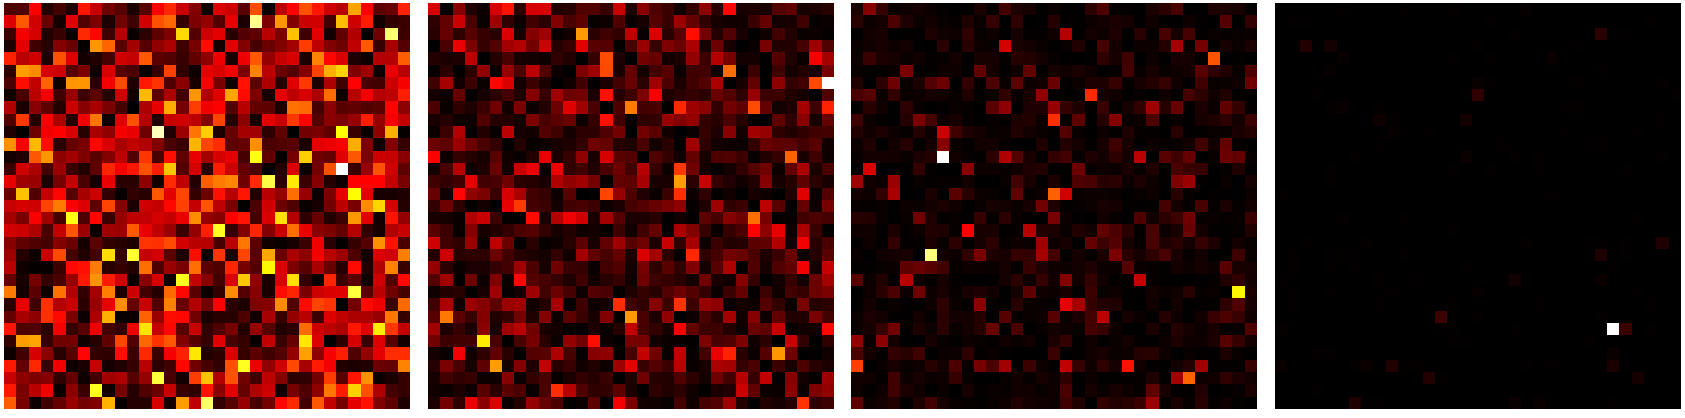

(a) ℓ_2 (b) ℓ_1 (c) $\ell_{1/2}$ (d) Student's/Cauchy

$$p_{prior}(u_i) \propto \exp(-|u_i|^p) \qquad vs. \qquad p_{prior}(u_i) \propto \frac{1}{1+u_i^2}$$

Why HBM? EEG/MEG Source Reconstruction

Aim: Reconstruction of brain activity by non-invasive measurement of induced electromagnetic fields (bioelectromagnetism) outside of the skull.

source: Wikimedia Commons

source: Wikimedia Commons

Why HBM? EEG/MEG Source Reconstruction

Aim: Reconstruction of brain activity by non-invasive measurement of induced electromagnetic fields (bioelectromagnetism) outside of the skull.

source: Wikimedia Commons

source: Wikimedia Commons

Notoriously ill-posed problem!

HBM for EEG/MEG Source Reconstruction

- Inversion with log-concave priors (e.g., ℓ_1-type) suffers from systematic depth miss-localization, HBM does not.
- HBM shows promising results for focal brain networks with simulated and real data and EEG-MEG combination.

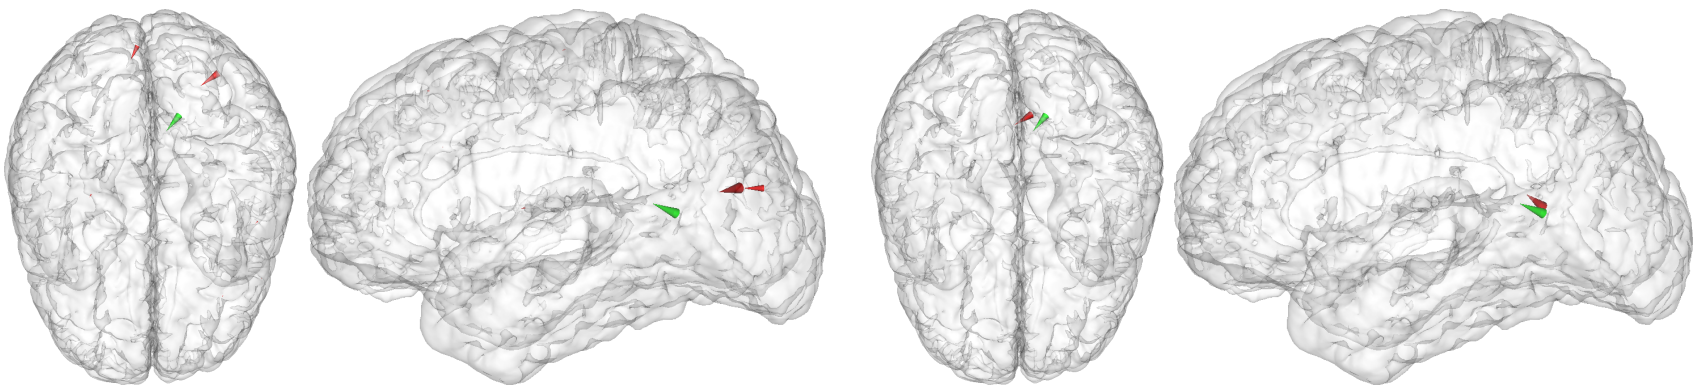

L., Pursiainen, Burger, Wolters, 2012. *Hierarchical Bayesian inference for the EEG inverse problem using realistic FE head models: Depth localization and source separation for focal primary currents.*
NeuroImage, 61(4):1364–1382.

Comparison: Two Approaches to Sparsity

feature	ℓ_p prior	HBM
$\mathcal{J}(u)$	$\|u\|_p^p$	$\frac{\nu+1}{2}\sum \log\left(1+\frac{u^2}{\nu\theta}\right)$
sparsifying parameter	$p > 0$	$\nu > 0$
quadratic limit	$p = 2$	$\nu \to \infty$
sparse limit	$p \to 0$	$\nu \to 0$
limit functional	$\|u\|_0$	$\sum_i^n \log(\|u_i\|)$ if all $u_i \neq 0$, $-\infty$ else
solutions	sparse	compressible
differentiable	$p > 1$	always
convex	everywhere for $p \geqslant 1$	$\|u\|_\infty < \sqrt{\nu\theta}$
homogeneous	yes	no

Comparison: Two Approaches to Sparsity

feature	ℓ_p prior	HBM
$\mathcal{J}(u)$	$\|u\|_p^p$	$\frac{\nu+1}{2}\sum \log\left(1+\frac{u^2}{\nu\theta}\right)$
sparsifying parameter	$p>0$	$\nu>0$
quadratic limit	$p=2$	$\nu\to\infty$
sparse limit	$p\to 0$	$\nu\to 0$
limit functional	$\|u\|_0$	$\sum_i^n \log(\|u_i\|)$ if all $u_i\neq 0$, $-\infty$ else
solutions	sparse	compressible
differentiable	$p>1$	always
convex	everywhere for $p\geqslant 1$	$\|u\|_\infty<\sqrt{\nu\theta}$
homogeneous	yes	no

Why not combine them to best (or worst?) of both worlds?

ℓ_p-hypermodels with generalized Gamma hyperpriors

$$p_{prior}(u,\gamma) \propto \exp\left(-\sum_i \left(\frac{|D_i^T u|^p}{\gamma_i} + \frac{\gamma_i^r}{\beta} - (r\alpha - 1 - 1/p)\log(\gamma_i)\right)\right)$$

Implicit prior with inverse gamma hyperprior:

$$\prod_i \left(1 + \frac{|D_i^T u|^p}{\beta}\right)^{-\alpha-1/p}$$

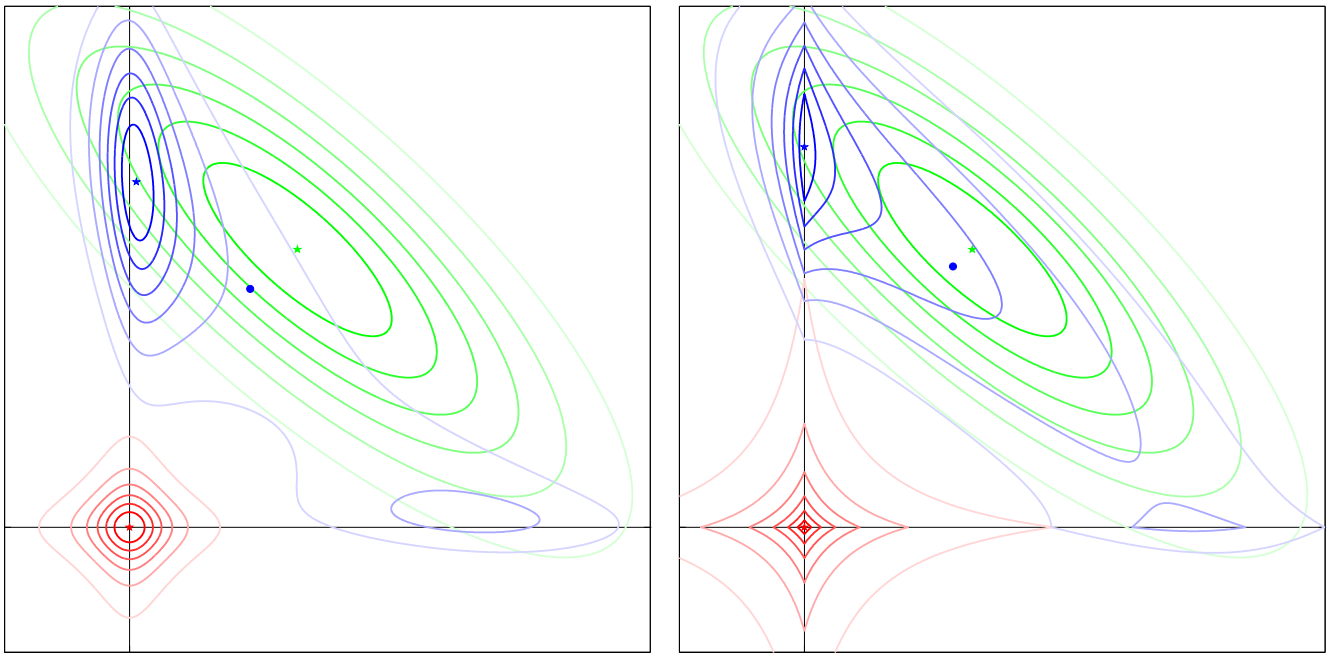

(a) $p = 2$ (b) $p = 1$

Posterior with gamma hyperprior ($r = 1$), $p = 1$, and $\alpha = 2$:

$$p_{post}(u|f) \propto \exp\left(-\frac{1}{2}\|f - A\,u\|_2^2 - \sum_i \left(\frac{|D_i^T u|}{\gamma_i} + \frac{\gamma_i}{\beta}\right)\right)$$

Computational scheme for full-MAP estimation equivalent to majorization-minimization scheme for $\ell_{1/2}$ regularization (Adaptive Lasso):

$$u^{(k)} = \underset{u}{\operatorname{argmin}} \left\{ \frac{1}{2}\|f - A\,u\|_{\Sigma_\varepsilon^{-1}}^2 + \frac{1}{\sqrt{\beta}} \sum_i \frac{|D_i^T u|}{\sqrt{|D_i^T u|^{(k-1)}}} \right\}$$

Bekhti, L, Salmon, Gramfort, 2017. *Revisiting majorization-minimization for non-convex sparse regression from a hierarchical Bayesian perspective: application to M/EEG inverse problem*, *almost submitted, I swear*.

Severely under-determined problems $f = Au$:

Many sparse solutions consistent with data!

- Log-concave priors erase this ambiguity and yield single result.
- Traditional UQ measure do not capture these aspects.
- Can we preserve but quantify and structure the ambiguity?

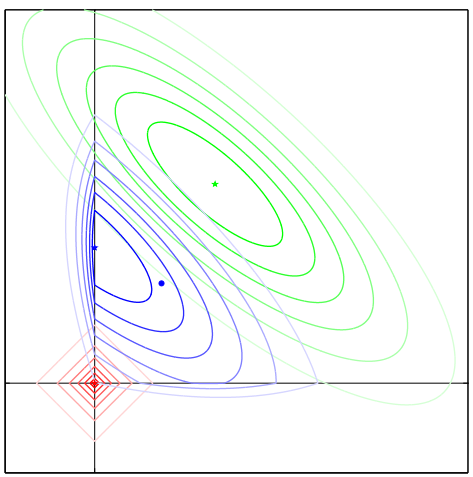

Severely under-determined problems $f = Au$:

Many sparse solutions consistent with data!

- Log-concave priors erase this ambiguity and yield single result.
- Traditional UQ measure do not capture these aspects.
- Can we preserve but quantify and structure the ambiguity?

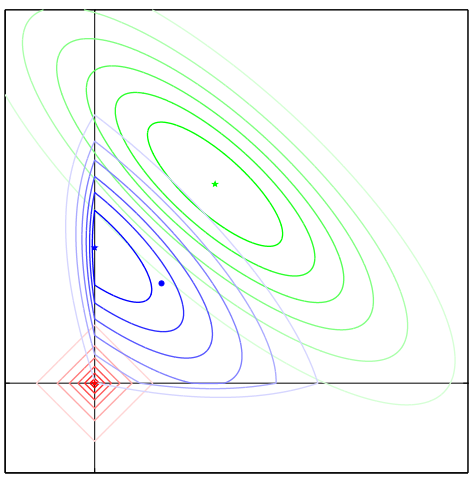

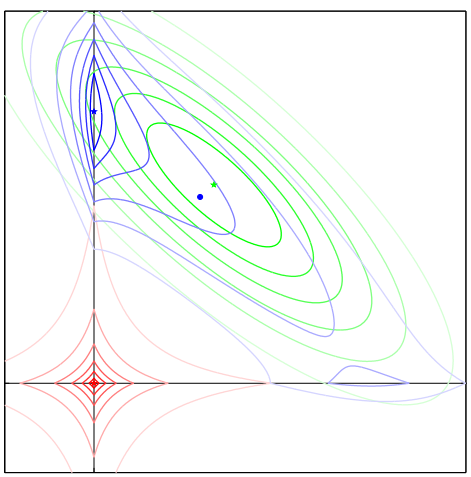

Each solution is a mode of the HBM posterior. Can we identify the most important ones, quantify their relative importance and visualize their structure?

Mode Analysis with MCMC & Optimization

- Generate MCMC chain of posterior samples.
- Use every sample as initialization of gradient-based optimization.
- Analyse resulting chain of modes.

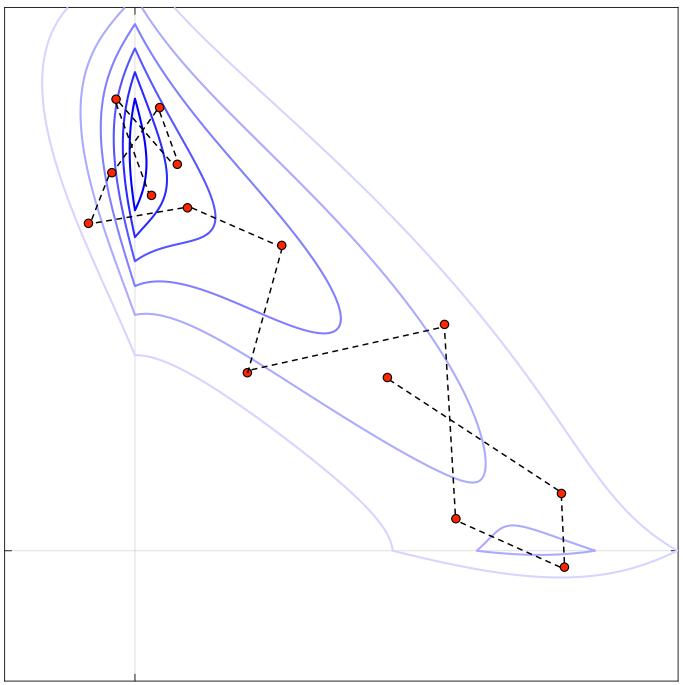

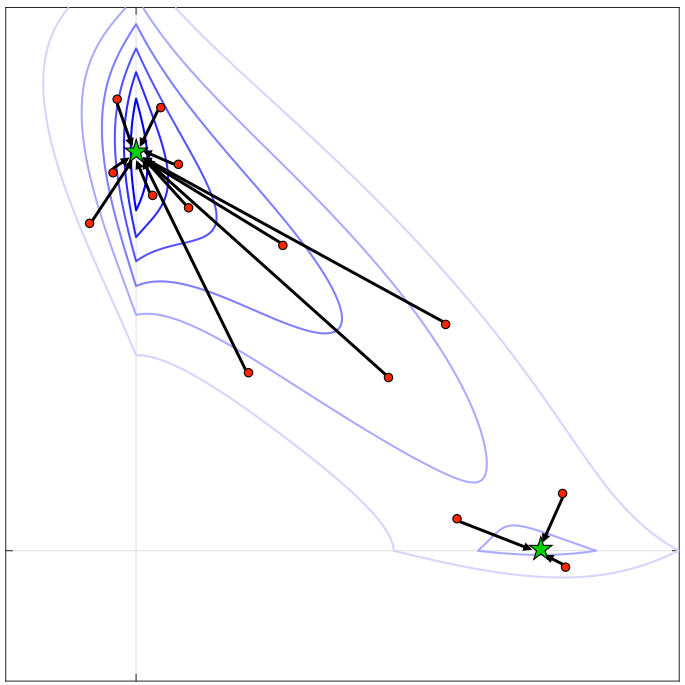

Sparse Source Network Analysis for EMEG Auditory Data

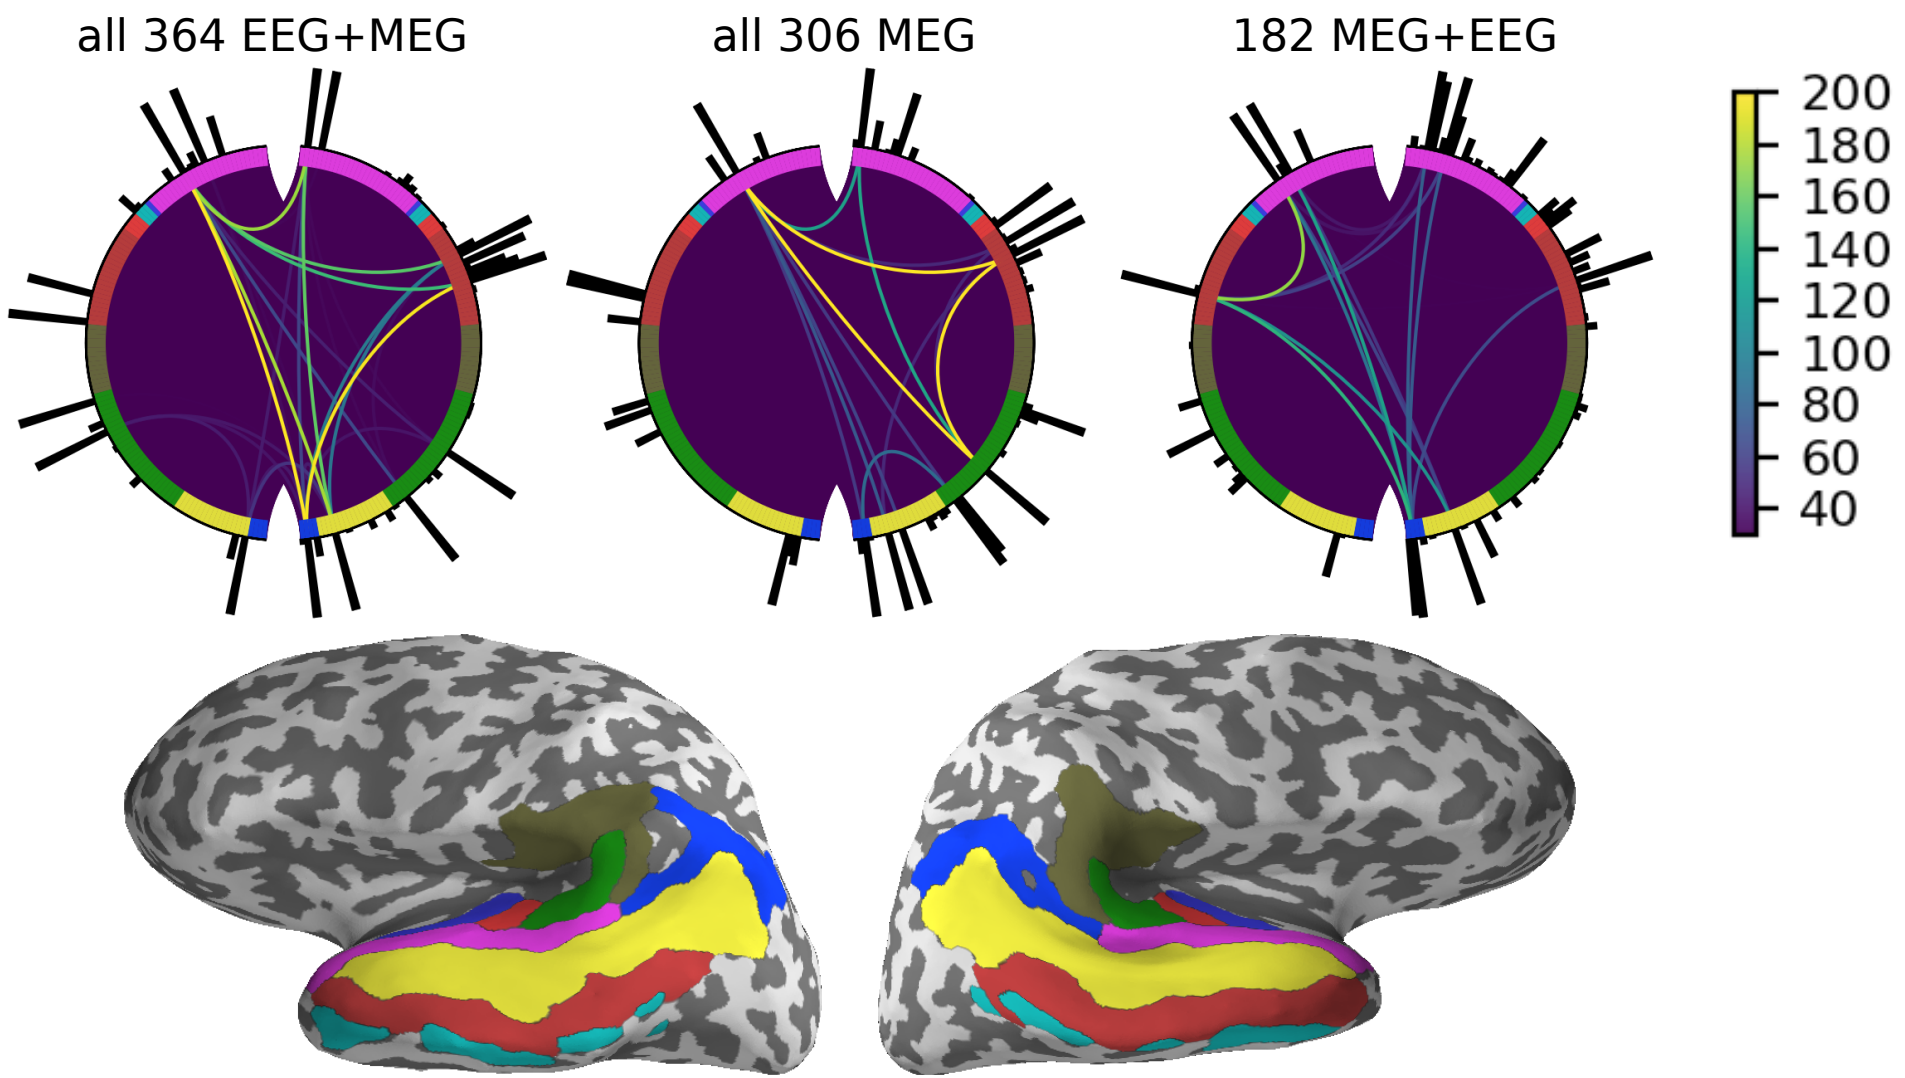

Summary, Outlook & Open Questions

- ℓ_p-norm and HBM road to sparsity: Neither perfect but (somewhat) computationally tractable. $\rightsquigarrow$ Slap and spike priors?
- MAP estimates are proper Bayes estimators, modes are meaningful.
- Meaningful UQ measures for sparse inversion / imaging that can complement variational approaches?
- Does it really make sense?
  (over confidence in ill-posed problems, prior domination)


Bekhti, L, Salmon, Gramfort, 2017. *Revisiting majorization-minimization for non-convex sparse regression from a hierarchical Bayesian perspective: application to M/EEG inverse problem*, *almost submitted, I swear!*.


L, 2016. *Fast Gibbs sampling for high-dimensional Bayesian inversion*, Inverse Problems.


L., 2014. *Bayesian Inversion in Biomedical Imaging*
PhD Thesis, University of Münster.


Burger, L, 2014. *Maximum-A-Posteriori Estimates in Linear Inverse Problems with Log-concave Priors are Proper Bayes Estimators*, Inverse Problems.


F, 2012. *Fast Markov chain Monte Carlo sampling for sparse Bayesian inference in high-dimensional inverse problems using L1-type priors*, Inverse Problems.


L, Pursiainen, Burger, Wolters, 2012. *Hierarchical Bayesian inference for the EEG inverse problem using realistic FE head models: Depth localization and source separation for focal primary currents*, *NeuroImage*.

Thank you for your attention!

Bekhti, L, Salmon, Gramfort, 2017. *Revisiting majorization-minimization for non-convex sparse regression from a hierarchical Bayesian perspective: application to M/EEG inverse problem*, *almost submitted, I swear!*.

L, 2016. *Fast Gibbs sampling for high-dimensional Bayesian inversion*, Inverse Problems.

L., 2014. *Bayesian Inversion in Biomedical Imaging*
PhD Thesis, University of Münster.

Burger, L, 2014. *Maximum-A-Posteriori Estimates in Linear Inverse Problems with Log-concave Priors are Proper Bayes Estimators*, Inverse Problems.

F, 2012. *Fast Markov chain Monte Carlo sampling for sparse Bayesian inference in high-dimensional inverse problems using L1-type priors*, Inverse Problems.

L, Pursiainen, Burger, Wolters, 2012. *Hierarchical Bayesian inference for the EEG inverse problem using realistic FE head models: Depth localization and source separation for focal primary currents*, *NeuroImage*.

MAP vs. CM Estimates: The Classical View

A theoretical argument "decides" the conflict: The Bayes cost formalism.

- An estimator is a random variable, as it relies on f and u.
- How does it perform on average? Which estimator is "best"?
- $\rightsquigarrow$ Define a cost function $\Psi(u, v)$.
- Bayes cost is the expected cost:

$$BC(\hat{u}) = \iint \Psi(u, \hat{u}(f))\, p_{like}(f|u)\, \mathrm{d}f\, p_{prior}(u)\, \mathrm{d}u$$

- Bayes estimator $\hat{u}_{BC}$ for given Ψ minimizes Bayes cost. Turns out:

$$\hat{u}_{BC}(f) = \underset{\hat{u}}{\operatorname{argmin}} \left\{ \int \Psi(u, \hat{u}(f))\, p_{post}(u|f)\, \mathrm{d}u \right\}$$

MAP vs. CM Estimates: The Classical View

Main classical arguments pro CM and contra MAP estimates:

- CM is Bayes estimator for $\Psi(u, \hat{u}) = \|u - \hat{u}\|_2^2$ (MSE).
- Also the minimum variance estimator.
- The mean value is intuitive, it is the "center of mass", the known "average".
- MAP estimate can be seen as an asymptotic Bayes estimator of

$$\Psi_\epsilon(u, \hat{u}) = \begin{cases} 0, & \text{if} \quad \|u - \hat{u}\|_\infty \leqslant \epsilon \\ 1 & \text{otherwise,} \end{cases}$$

  for $\epsilon \to 0$ (uniform cost). $\Longrightarrow$ It is not a proper Bayes estimator.

- MAP and CM seem theoretically and computationally fundamentally different $\Longrightarrow$ one should decide.
- *"A real Bayesian would not use the MAP estimate"*
- People feel "ashamed" when they have to compute MAP estimates (even when their results are good).

A False Conclusion

"A real Bayesian would not use the MAP estimate as it is not a proper Bayes estimator".

"MAP estimate can be seen as an asymptotic Bayes estimator of

$$\Psi_\epsilon(u, \hat{u}) = \begin{cases} 0, & \text{if} \quad \|u - \hat{u}\|_\infty < \epsilon \\ 1 & \text{otherwise,} \end{cases}$$

for $\epsilon \to 0$.
???$\Longrightarrow$??? It is not a proper Bayes estimator."

"MAP estimator is asymptotic Bayes estimator for some degenerate Ψ"
$\not\Rightarrow$ "MAP can't be Bayes estimator for some proper Ψ" !!!!

Two New Bayes Cost Functions

Define

(a) $\Psi_{\text{LS}}(u, \hat{u}) := \|A(\hat{u} - u)\|^2_{\Sigma_\varepsilon^{-1}} + \beta\|L(\hat{u} - u)\|_2^2$

(b) $\Psi_{\text{Brg}}(u, \hat{u}) := \|A(\hat{u} - u)\|^2_{\Sigma_\varepsilon^{-1}} + \lambda D_{\mathcal{J}}(\hat{u}, u)$

for a regular L and $\beta > 0$.

Properties:

- Proper, convex cost functions
- For $\mathcal{J}(u) = \beta/\lambda\|Lu\|_2^2$ (Gaussian case!) we have $\lambda D_{\mathcal{J}}(\hat{u}, u) = \beta\|L(\hat{u} - u)\|_2^2$, and $\Psi_{\text{LS}}(u, \hat{u}) = \Psi_{\text{Brg}}(u, \hat{u})$!

Theorems:

(I) The CM estimate is the Bayes estimator for $\Psi_{\text{LS}}(u, \hat{u})$

(II) The MAP estimate is the Bayes estimator for $\Psi_{\text{Brg}}(u, \hat{u})$

Bregman distances

For a proper, convex functional $\Psi : \mathbb{R}^n \longrightarrow \mathbb{R} \cup \{\infty\}$, the *Bregman distance* $D_\Psi^p(f, g)$ between $f, g \in \mathbb{R}^n$ for a subgradient $p \in \partial\Psi(g)$ is defined as

$$D_\Psi^p(f, g) = \Psi(f) - \Psi(g) - \langle p, f - g \rangle, \qquad p \in \partial\Psi(g)$$

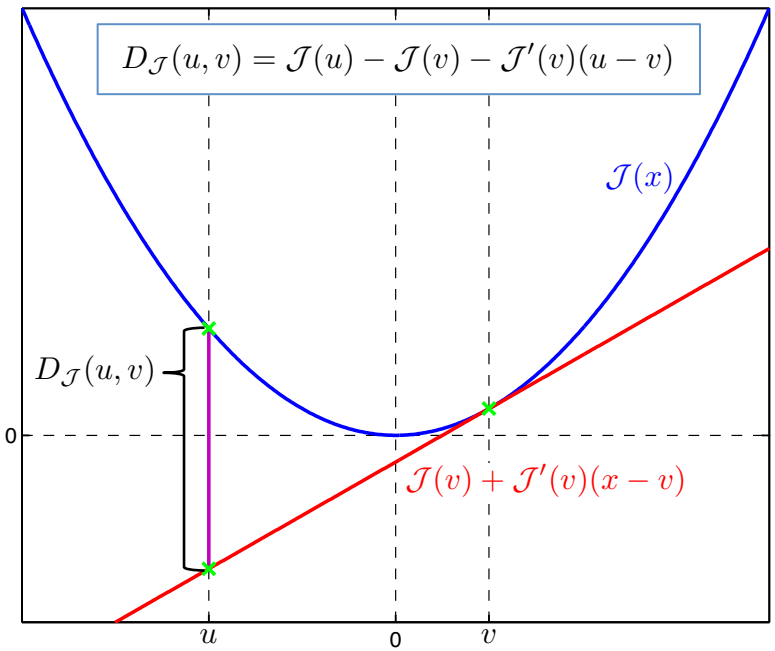

(c) $\mathcal{J}(x) = x^2$

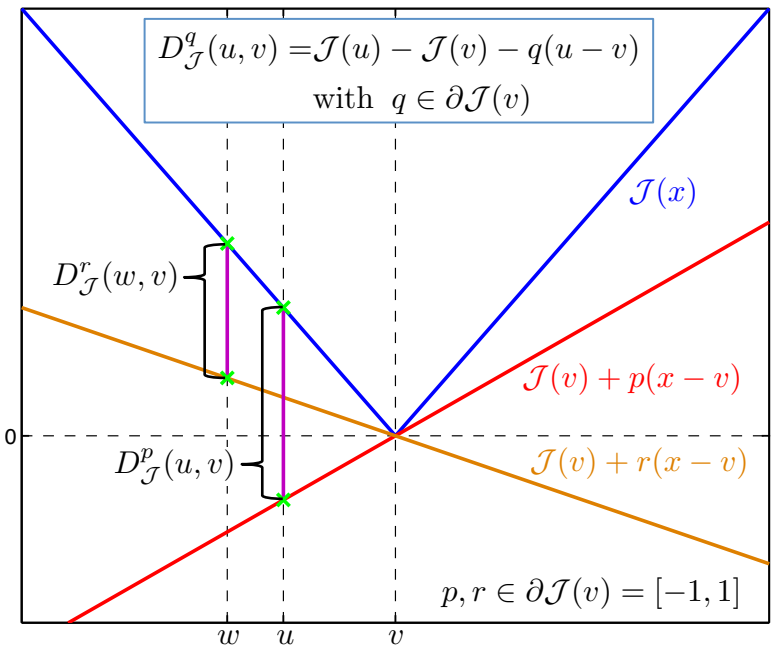

(d) $\mathcal{J}(x) = |x|$

Basically, $D_\Psi(f, g)$ measures the difference between Ψ and its linearization in f at another point g